

# CPSC 340 Assignment 6 (due Friday April 7 at 11:59pm)

## Neural Networks

### Instructions

Rubric: {mechanics:3}

The above points are allocated for following the general homework instructions on GitHub.

Note: this assignment is a bit more open-ended than the previous ones. There aren't really "correct answers". Try to have fun with it and learn something along the way. We'll mark generously... no need to go overboard!

### 1 Neural Networks


Rubric: {reasoning:5}

If you run `python main.py -q 1` it will train a neural network on the *basisData.pkl* data from Assignment 3. However, in its current form it uses no hidden layers and therefore is just performing linear regression. In fact, if you take a look at the figure that is generated, you should see that it matches the figure from Assignment 3 Question 2.1. Try to improve the performance of the neural network by adding hidden layers and tuning the other hyperparameters. **Hand in your plot after changing the code to have better performance, and list the changes you made. Write a couple sentences explaining why you think your changes improved the performance. When appropriate, refer to concepts from the course like overfitting or optimization.**

To see the list of hyperparameters and their definitions, see the scikit-learn MLPRegressor documentation: [http://scikit-learn.org/stable/modules/generated/sklearn.neural\\_network.MLPRegressor.html](http://scikit-learn.org/stable/modules/generated/sklearn.neural_network.MLPRegressor.html).

Note: "MLP" stands for Multi-Layer Perceptron, which is another name for artificial neural network.

Also: we didn't discuss L-BFGS in class, but it may help to set `solver` to `lbfgs` since the data set is tiny.



I changed the Regressor to 100 neurons in the 1st hidden layer, and added an alpha of 0.01 to regularize the error, and optimize the hyper-parameters, which in turn reduces overfitting. Came to the number 100 after trying various numbers between 10-1000 and tried adding layers as well, but came to the conclusion only one hidden layer was present.

## 2 Your Own Data

Rubric: {reasoning:5}

Try a neural network on the data set of your choosing. You can use a data set that's built into scikit-learn (<http://scikit-learn.org/stable/datasets/>), or you can use a data set from a previous homework assignment, or something from an area that interests you. If you are doing classification rather than regression, use scikit-learn's MLPClassifier rather than the MLPRegressor that we've provided in the assignment code: [http://scikit-learn.org/stable/modules/generated/sklearn.neural\\_network.MLPClassifier.html](http://scikit-learn.org/stable/modules/generated/sklearn.neural_network.MLPClassifier.html)

Try fiddling with the hyperparameters. Write a few sentences about what you tried, what train/validation error you achieved, and whether you think neural networks are a good model for your dataset. Optionally (for glory rather than marks), install Keras/TensorFlow and compare it with the scikit-learn neural net.

For the Iris dataset in scikit, on using the MLPClassifier with no hidden layers gives Training error = 0.0 and Validation error = 0.6, while increasing the number of layers doesn't really decrease the validation error, on repeated runs, with and without hidden layers give very similar errors.